

# SOCIAL NETWORK ANALYSIS for DATA SCIENTISTS

today's menu: LECTURE: ERGM III (LECTURE Week 7)

Your lecturer: Claudia & Gergo

Playdate: October, 15th, 2025

# Menu'



- ERGMs Estimation
- Model diagnostics
- Curved Terms (to solve degeneracy problem)
- Goodness of fit

# *ERGMs Estimation*

# Estimation



The estimation allows for the identification of the parameters that maximize the likelihood of a graph.

Finding the  $\theta$  of the terms that maximize the probability of simulating graphs like the observed real one

- We discussed how to insert terms and how to specify the models
- We still need to check if the specification we inserted is doing its job
- Since the algorithm takes care of the estimation we need to understand it a bit better

# ERGM algorithms



**Dyadic independent ERGM:** estimated as a logit model with Maximum (pseudo) Likelihood Estimation (MLE or MPLE)

- We need to check the goodness of fit as we did with logit models
- We do that using the `ergm:::gof()` function

**Dyadic (inter)dependent:** we need to estimate parameters using Markov chains Monte Carlo MCMC simulations

- We still need to check the goodness of fit but also we need to check whether the simulation worked as we wanted to. We need to check two indicators
  - `ergm:::gof()`
  - `ergm:::mcmc.diagnostics()`

# *Intuitive estimation -Dyadic dependent*



## Explaining a Markov Chain Maximum Likelihood estimation

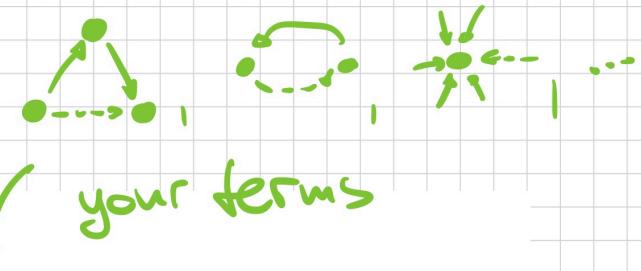
- We simulate Markov Chains using a Monte Carlo (random) method.
- A Markov chain is a sequence of random variables such that the value taken by the random variable only depends upon the value taken by the previous variable.
- By switching the entries in the matrix representing the network (0-1) we generate a series of graphs that differ from each other only by one edge and that depend from the previous graph.
- This sequence of generated graphs is a Markov chain.
- Within this generation of chains the Metropolis-Hastings algorithm or the Gibbs sampler decide which entries (0-1) change state.
- These algorithms start picking a variable at random and change its state
- They retain the change of state only if the new graph increases the likelihood
- Finally the Robbins-Monroe algorithm estimates the parameters in the model.

ERGM

MATH

$$p(X = x) = \frac{\exp \{\theta \cdot v(G)\}}{\exp \left\{ \log \left( \sum_{y \in Y} \exp \{\theta \cdot v(G)\} \right) \right\}}$$

$$\exp \{ \theta \cdot v(G) \}$$



$$\exp \{ \theta \cdot v(G) \}$$

"strength" of  
your terms



your terms

What are your terms?

Remember lab ERGM 1!

(Markovian dependence, neighbourhood, etc)

probability of your  
specific network forming



$$p(X = x) = \frac{\exp\{\theta \cdot v(G)\}}{\exp\left\{\log\left(\sum_{y \in Y} \exp\{\theta \cdot v(G)\}\right)\right\}}$$

probability of your  
specific network forming

(numerical representation of)

your observed  
network

$$p(X = x) = \frac{\exp \{\theta \cdot v(G)\}}{\exp \left\{ \log \left( \sum_{y \in Y} \exp \{\theta \cdot v(G)\} \right) \right\}}$$

probability of your  
specific network forming

(numerical representation of)

your observed  
network

$$p(X = x) = \frac{\exp \{ \theta \cdot v(G) \}}{\exp \left\{ \log \left( \sum_{y \in Y} \exp \{ \theta \cdot v(G) \} \right) \right\}}$$

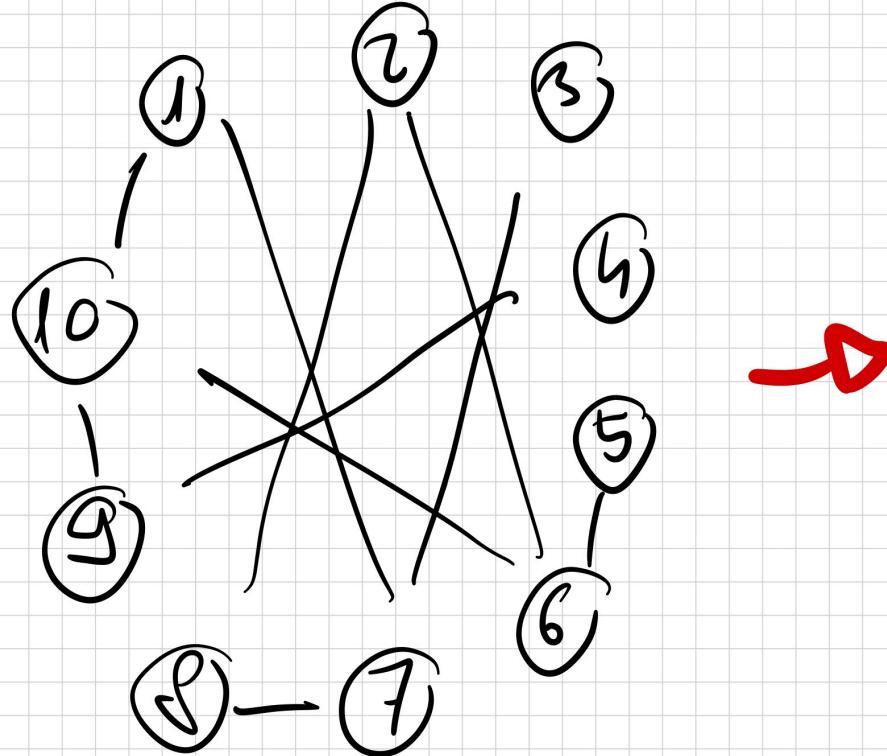
(numerical representation of) all possible networks

# PROBLEM!

"all possible networks" is an

insanely large number of networks.

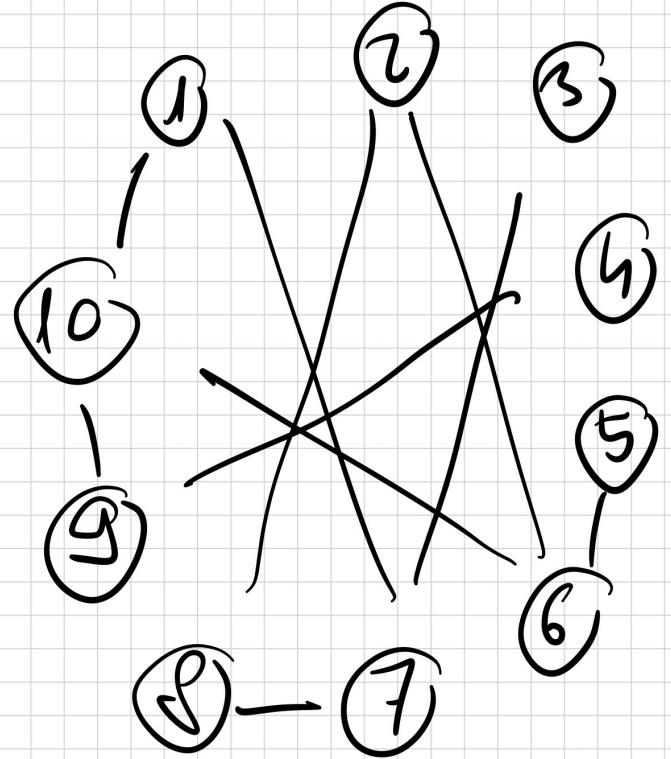
For nerds:  $n$  nodes define  $2^{\frac{n(n-1)}{2}}$  graphs.



35186372088832

different graphs  
are possible





35 186 372 088 832

different graphs  
are possible



same magnitude as :

- number of ants alive on Earth
- population of the Star Wars and Asimov universes

( - more than how many cells )  
you have

TLDR :

Computing an ERGN is  
impossible.

( except in case of dyadic independence ,

!!

network doesn't matter

$$p(X = x) = \frac{\exp\{\theta \cdot v(\text{red})\}}{\exp\left\{\log\left(\sum_{y \in Y} \exp\{\theta \cdot v(\text{red})\}\right)\right\}}$$

no network = no  $10^{15}$  combinations ✓ )

Unless you know a clever trick:

MCMC

Markov Chain Monte Carlo

# Goal :

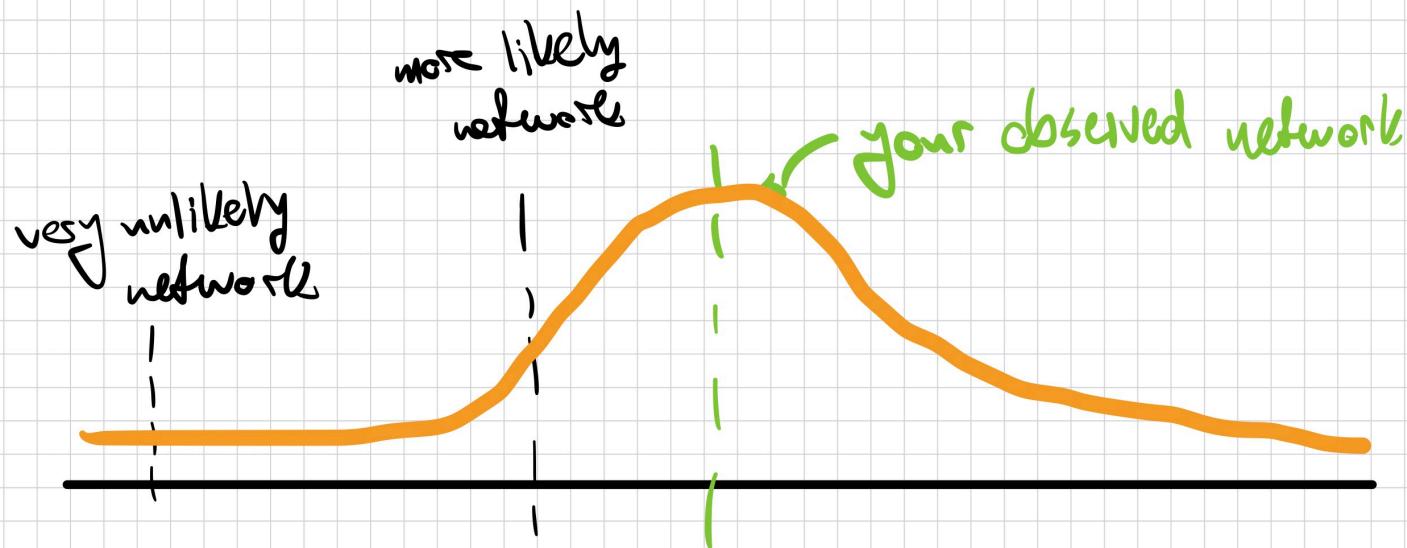
Identify  $\mathcal{W}$  that maximizes probability of observed graph.



this is what you can interpret to examine hypotheses.

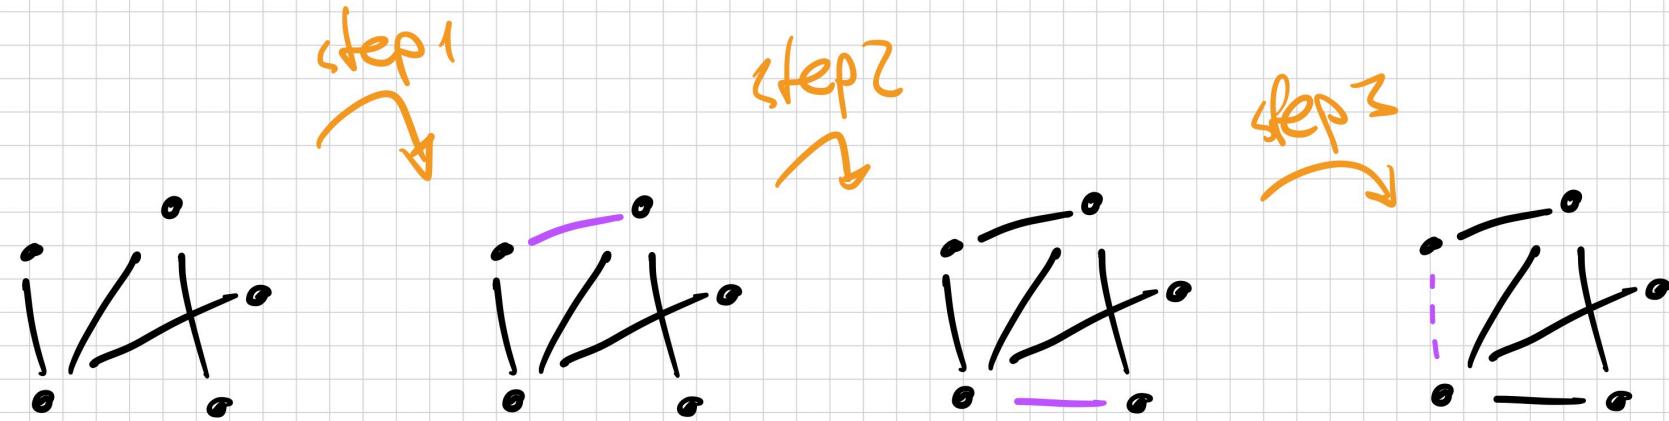
# Need :

Sample of networks from the same probability distribution. (described by  $\Theta$ )



Let's  
simulate!

Markov-chaine = series of networks, edge changes

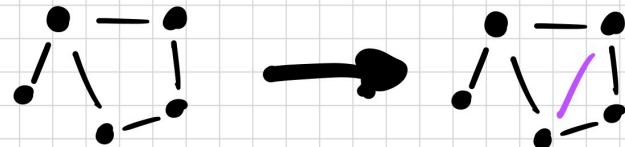


This leads to a sample of networks, but we  
don't know the distribution (random change)

# HOW TO PICK WHICH EDGE TO CHANGE?

## METROPOLIS - HASTINGS:

1. change edge at random

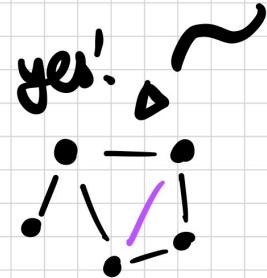


2. Assess if "good change"

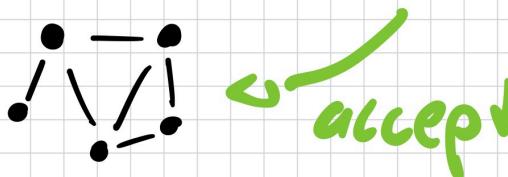
yes = changed network is base

$$P(\cdot \overset{?}{\sim} \cdot) > P(\cdot \sim \cdot)$$

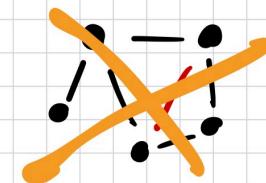
no = old network is base



3. Rinse and repeat!



nope!

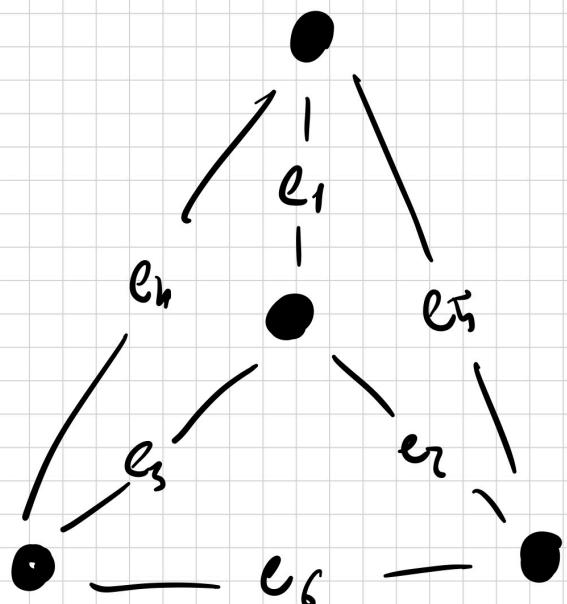


reject

## ALTERNATIVE WAY:

Gibbs - sampler

→ making sure the steps go in the right direction



1. calculate  $P(e_1 | \text{rest of net})$

↳ P small? → set  $e_1$  to  $\emptyset$

2. calculate  $P(e_2 | \text{rest of net without } e_1)$

↳ P big → set  $e_2$  to 1

3 ... rinse and repeat

# HOW DOES THIS GET US A BETTER $\Theta$ ?

Robbins-Monro : Example : getting a better  $\Theta$  value for triadic closure

1. Observed network has 4 closed triads
2. Using current  $\Theta_t$  and MCMC, you make a sample network
3. Count number of triads in sample : 3
4. New  $\Theta_{t+1} = \Theta_t + \alpha_t (4 - 3)$  ← here is where you make it better  
  
 $\alpha$  is a learning rate

# Convergence



- Chains should evolve toward increasing the likelihood of producing networks similar to the observed one.
- If the process goes according to plans, we say that the model **converged**
- If you look at the simulation output you will see a series of lines with different trends that chain after chain become similar and overlap (We will see it in a few minutes)
- If the model does **not converge** the chains never become similar and parallel universe continue to exist until the model fails (We will also see that in a few minutes)

# Burn-in



- In the beginning, the model will have a time for the Metropolis-Hastings algorithm or the Gibbs sampler to increase the likelihood of the chains to resemble the real network
- This is the **burn-in**
- The burn-in represents the number of simulations that are removed from the MCMC in order to make the chain "forget" the (random) starting point

# MCMC sample size



- If the model struggles at convergence it is a good idea to give it the opportunity to be more accurate.
- Increasing the sample size, the Metropolis-Hastings algorithm or the Gibbs sampler can source more information about the previous chain and increase the accuracy of the following one.
- It does not always work. If your terms are not capturing the dynamics of the real network it will never work.
- It increases your computational time, so careful when you use it

# *Model diagnostics*

# Checking the simulation in Dyadic dependent Models



## Model using Florentine Business network

```
fit <- ergm:::ergm(floB ~ edges + degree(1))  
ergm:::mcmc.diagnostics(fit)
```

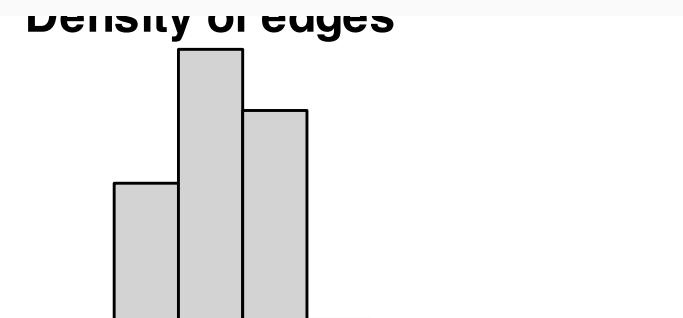
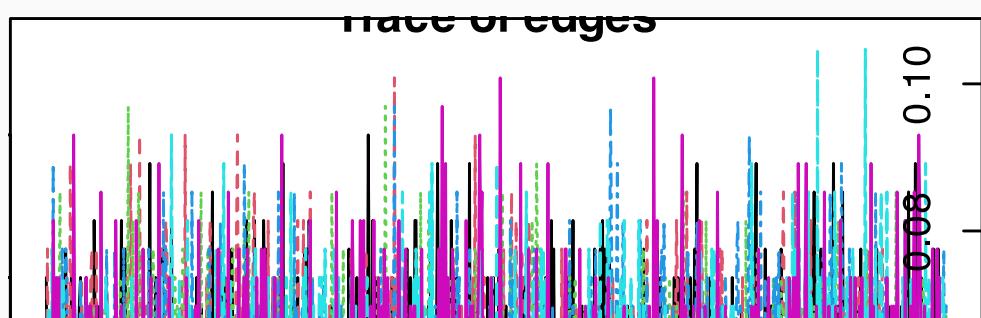
This function outputs three things:

- Text that explains the diagnostics
- A series of relevant measures
- Plots the MCMC diagnostics

# MCMC Diagnostics



```
fitFlo.01 <- ergm:::ergm(floB ~ edges + degree(1),  
                           control = ergm:::control.ergm(  
                             MCMC.samplesize = 5000,  
                             MCMC.burnin = 1000,  
                             MCMLE.maxit = 10,  
                             parallel = 6,  
                             parallel.type = "PSOCK"  
                           ))  
  
ergm:::mcmc.diagnostics(fitFlo.01)
```



# Understanding the `mcmc.diagnostics` output



You need to know that

- your diagnostic analyses the iterations of the last chain produced
- ERGMs, as OLS, is predicting mean values of parameters
- The mean in the diagnostics should be 0
- the standard deviation should form a bell curve

"Quantiles for each variable": 2.5% 25% 50% 75% 97.5%

-8, -2, 0, 2, 8 Would be an ideal scenario

"Are sample statistics significantly different from observed?"

You want your p-value here to be as high as possible

# Plot of the MCMC diagnostics

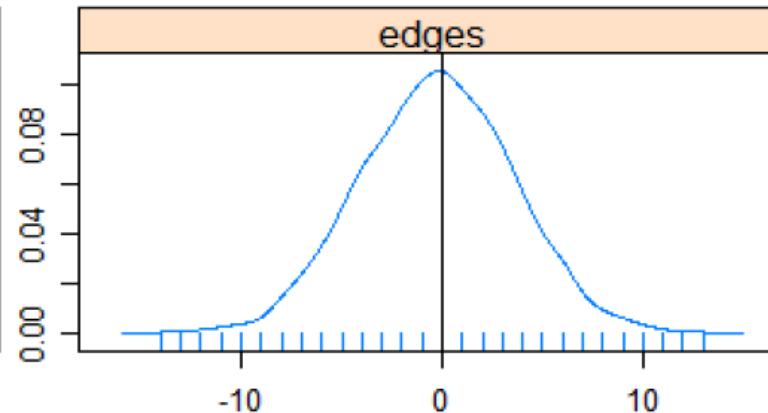
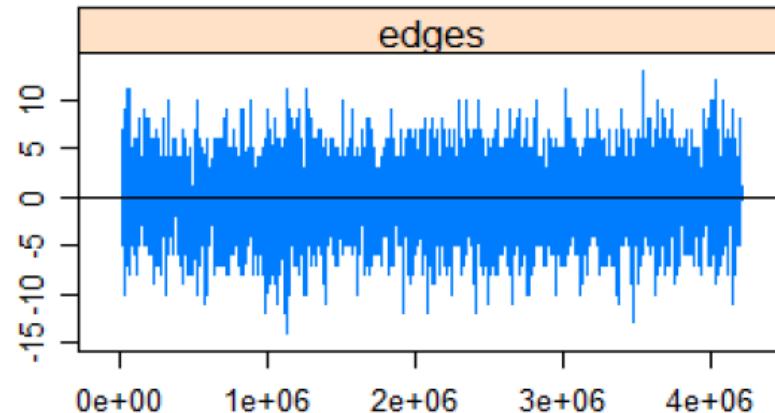
DS

statistics from the final iteration of the MCMC chain

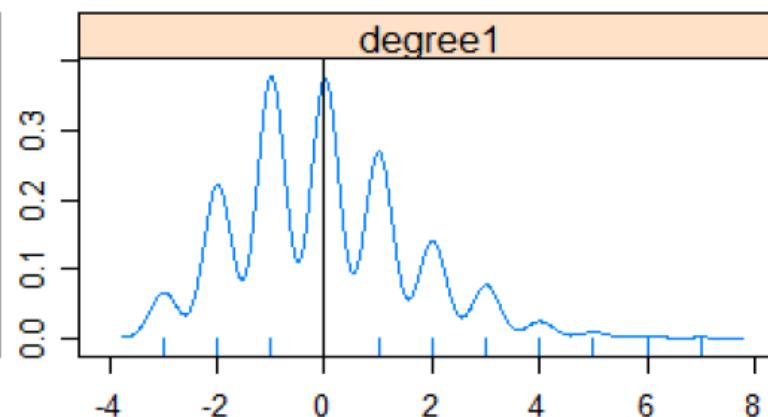
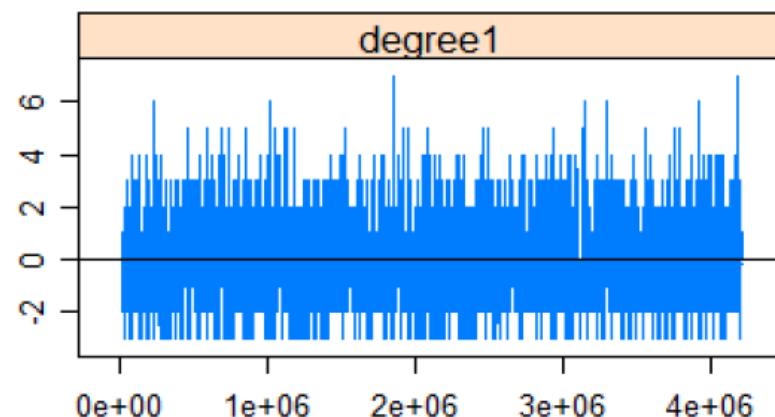
THIS IS GOOD!

Sample statistics

"traceplot"  
(the deviation  
of the statistic  
value in each  
sampled  
network from  
the observed  
value)



distribution  
of the sample  
statistic  
deviations  
should  
approximate  
a bell curve



This ok since  
this is a small  
network

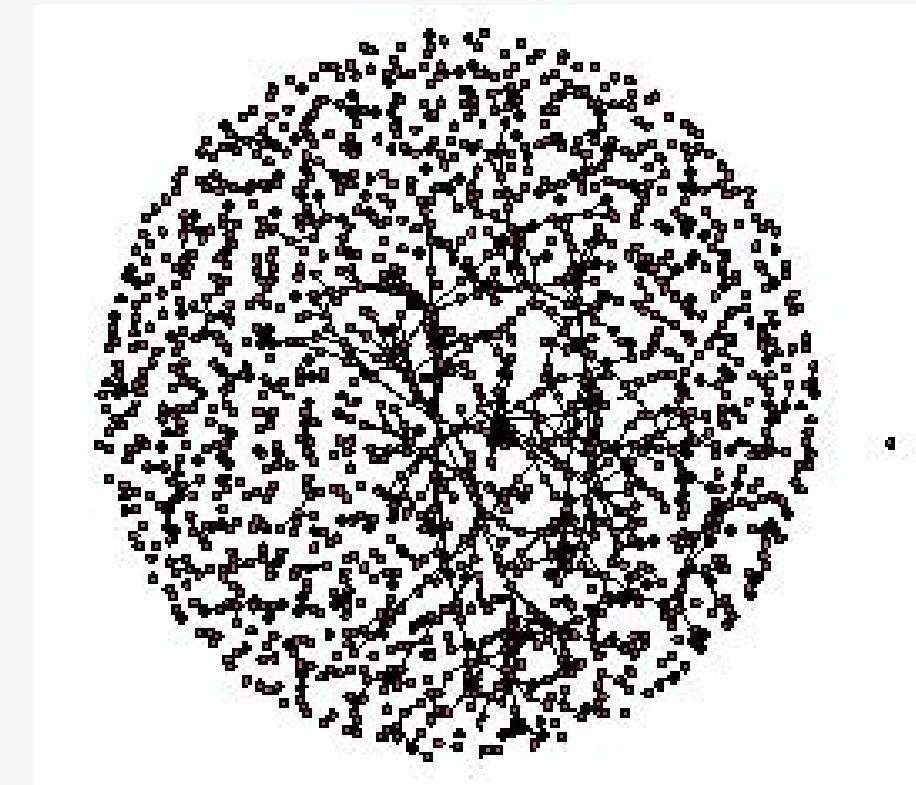
# Trace Plot shape: fuzzy caterpillar



# *Bad model example*

Faux Magnolia High. Data simulated from an American High School

- Nodes: 1461
- Edges: 974
- Sparse
- Undirected



# Bad model example



Running this model we will get an error

```
# fitMag.01 <- ergm::ergm(magnolia ~ edges + triangles,  
#                           control = ergm::control.ergm(  
#                                         MCMC.samplesize = 5000,  
#                                         MCMC.burnin = 1000,  
#                                         MCMLE.maxit = 10,  
#                                         parallel = 6,  
#                                         parallel.type = "PSOCK"  
# ))
```

Error in ergm.MCMLE(init, nw, model, initialfit = (initialfit <- NULL), :

Number of edges in a simulated network exceeds that in the observed by a factor of more than 20. This is a strong indicator of model degeneracy or a very poor starting parameter configuration. If you are reasonably certain that neither of these is the case, increase the MCMLE.density.guard control.ergm() parameter.

This is called model DEGENERACY

# The term triangles caused Degeneracy



Instead of converging, the algorithm heads off into networks that are much much more dense than the observed network. This is such a clear indicator of a degenerate model specification that the algorithm stops after 3 iterations, to avoid heading off into areas that would cause memory issues.

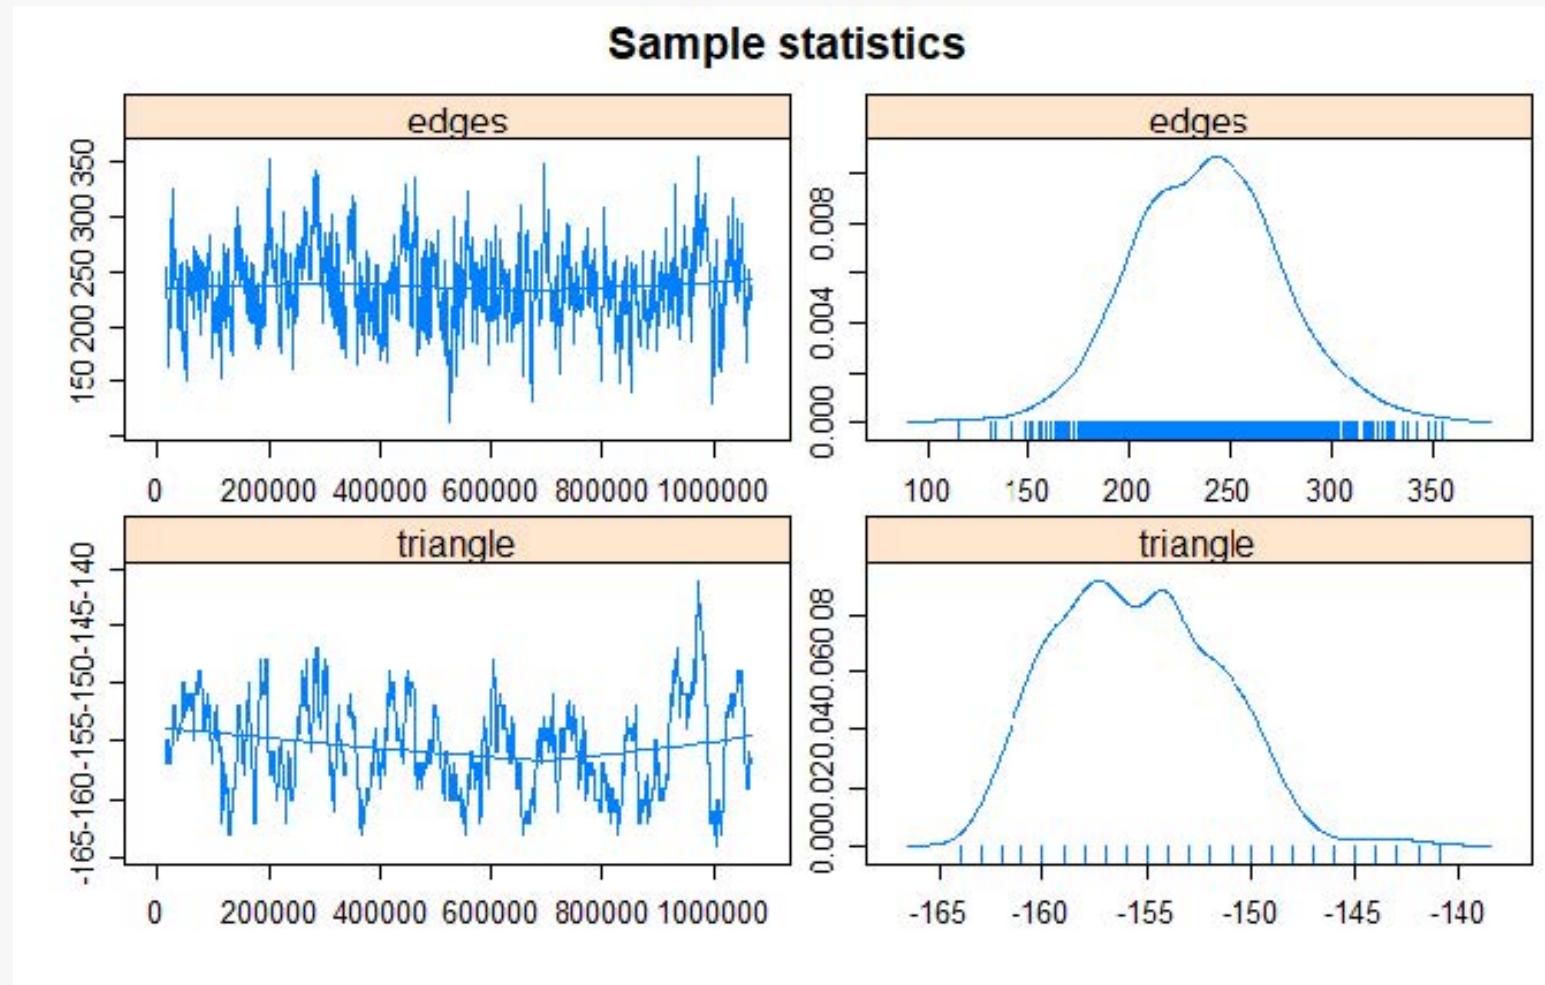
We can run the model limiting the number of iterations so it will converge and we can see what are the problems  
Since the 3rd iteration was the last one in the model output we stop it at 1

```
# fitMag.01 <- ergm::ergm(magnolia ~ edges + triangles,  
#                           control = ergm::control.ergm(  
#                                         MCMC.samplesize = 5000,  
#                                         MCMC.burnin = 1000,  
#                                         MCMLE.maxit = 1,  
#                                         parallel = 6,  
#                                         parallel.type = "PSOCK")  
#  
#  
# ergm::mcmc.diagnostics(fitMag.01)
```

Since sometimes this does not converge I'm not running the code in the slides

# Bad model trace plot (where is the 0?!)

Chains are not mixing well!

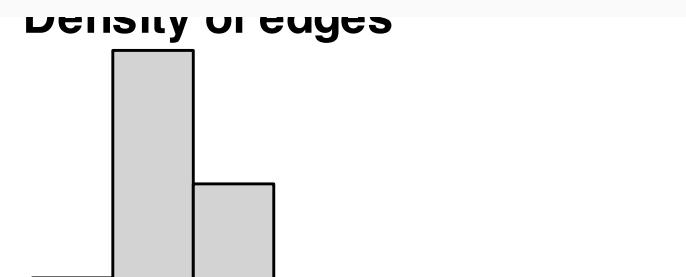
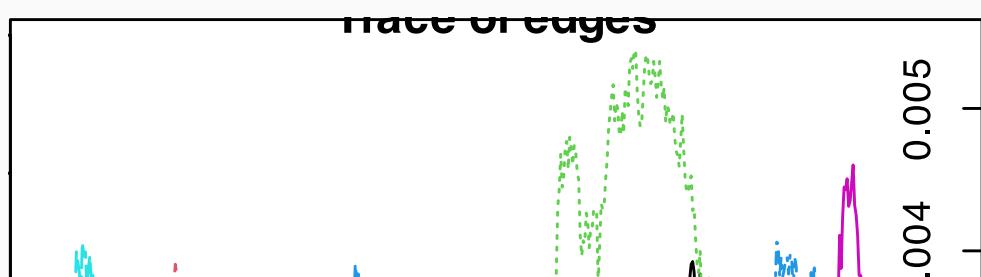


*Curved Terms (to solve degeneracy problem)*

# How do we make the model converge?

We substitute the Markovian term with a curved one

```
set.seed(000023)
Magfit.02 <- ergm:::ergm(magnolia ~ edges + gwesp(0.25, fixed = TRUE),
                           control = ergm:::control.ergm(
                               MCMC.samplesize = 5000,
                               MCMC.burnin = 1000,
                               MCMLE.maxit = 1,
                               parallel = 6,
                               parallel.type = "PSOCK")
                           )
ergm:::mcmc.diagnostics(Magfit.02)
```



# Curved terms



- do the same job as the Markovian terms
- but they have been mathematically improved to avoid degeneracy and improve model fit

When you use a Markovian term that does not work swap it for the curved equivalent.

- DISCLAIMER: Markovian terms will work only for very small and simple networks.

Keep your Curved equivalent handy.

# Most common Curved terms

- `gwesp(decay = 0.25, fixed = FALSE)` - it can be used in place of `triangles` or `esp`
- `gwdsp(decay = 1, fixed = FALSE)` - it can be used in place of `dsp`
- `gwnsp(decay = 1, fixed = FALSE)` - it can be used in place of `nsp`
- `gwdegree(decay = 0.5, fixed = FALSE)` - it can be used in place of `degree`, but careful, this is degree distribution.
- `gwidegree(decay = 0.5, fixed = FALSE)` - it can be used in place of `idegree`, but careful, this is degree distribution.
- `gwodegree(decay = 0.5, fixed = FALSE)` - it can be used in place of `odegree`, but careful, this is degree distribution.
- there are also curved terms that don't have a direct correspondence to the Markovian one. E.g.,  
`dgwdsp(decay, fixed=FALSE` - it can be used as a sum of `gwnsp` and `gwesp`

The optional argument `fixed` indicates whether the decay parameter is fixed at the given value, or is to be fit as a curved exponential family model. AKA: it uses different algorithms. At this stage, try both and keep the best one.

# How to find curved terms?



```
ergm::search.ergmTerms(search = "curved")

## Found 16 matching ergm terms:
## altkstar(lambda, fixed=FALSE) (binary)
##     Alternating k-star
##
## Curve(formula, params, map, gradient=NULL, minpar=-Inf, maxpar=+Inf, cov=NULL) (binary, valued)
## Parametrise(formula, params, map, gradient=NULL, minpar=-Inf, maxpar=+Inf, cov=NULL) (binary, valued)
## Parametrize(formula, params, map, gradient=NULL, minpar=-Inf, maxpar=+Inf, cov=NULL) (binary, valued)
##     Impose a curved structure on term parameters
##
## gwb1degree(decay, fixed=FALSE, attr=NULL, cutoff=30, levels=NULL) (binary)
##     Geometrically weighted degree distribution for the first mode in a bipartite network
##
## gwb1dsp(decay=0, fixed=FALSE, cutoff=30) (binary)
##     Geometrically weighted dyadwise shared partner distribution for dyads in the first bipartition
##
## gwb2degree(decay, fixed=FALSE, attr=NULL, cutoff=30, levels=NULL) (binary)
##     Geometrically weighted degree distribution for the second mode in a bipartite network
```

# *Goodness of fit*

Let's get back to our pretty Florentine Business network

```
FloFit.01 <- ergm::ergm(floB ~ edges + degree(1))
fit.gof<- ergm::gof(FloFit.01)
fit.gof

##
## Goodness-of-fit for degree
##
##      obs min mean max MC p-value
## degree0   5  0 2.82   7     0.32
## degree1   3  0 3.15   8     1.00
## degree2   2  0 5.51  10     0.08
## degree3   2  0 2.92   8     0.88
## degree4   3  0 1.34   5     0.30
## degree5   1  0 0.19   3     0.28
## degree6   0  0 0.07   1     1.00
##
## Goodness-of-fit for edgewise shared partner
##
##      obs min mean max MC p-value
```

# GOF output



## Output: 4 Parts

1. Goodness of fit for degree (always)
2. Goodness of fit for edgewise shared partner (always)
3. Goodness of fit for minimum geodesic distance (always)
4. Goodness of fit for model statistics (changes with model specifications)

It is possible to customize some other parameters

# GOF 1. Goodness-of-fit for degree

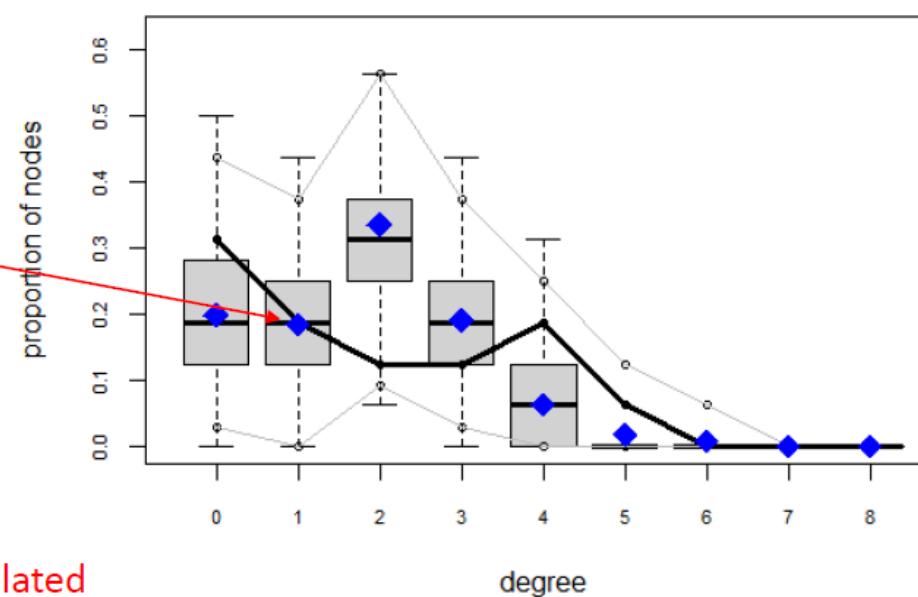
```
fit <- ergm:::ergm(flobusiness ~ edges + degree(1))
fit.gof <- ergm:::gof(fit)
plot(fit.gof)
```

Good! Real network == Network created with model parameters!

Goodness-of-fit for degree

obs	min	mean	max	MC	p-value
0	5	0	3.19	9	0.38
1	3	1	3.50	7	1.00
2	2	2	4.98	9	0.12
3	2	0	2.92	7	0.88
4	3	0	1.07	4	0.30
5	1	0	0.28	3	0.48
6	0	0	0.06	1	1.00

One row for each possible node degree



simulated

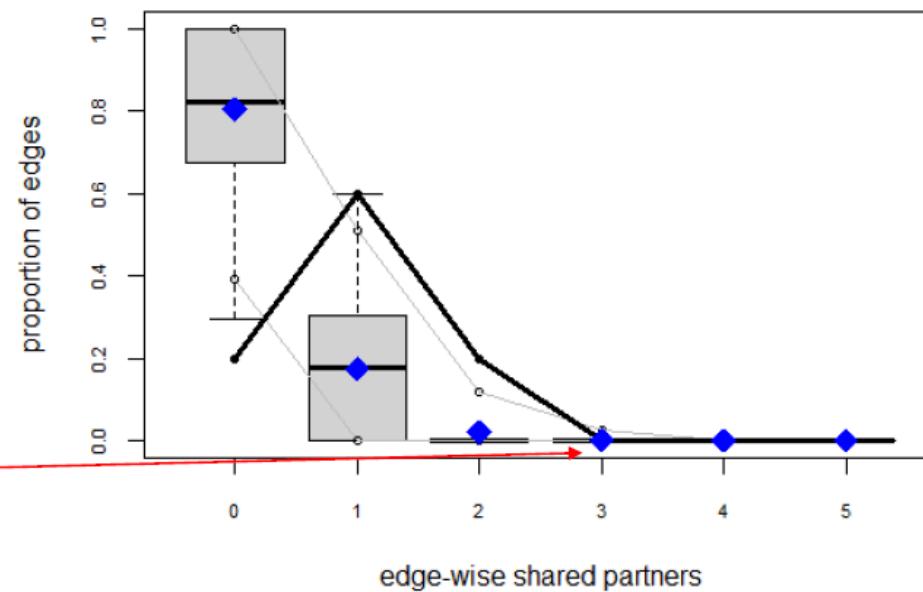
## GOF 2. Goodness-of-fit for edgewise shared partner

```
fit <- ergm::ergm(flobusiness ~ edges + degree(1))  
fit.gof <- ergm::gof(fit)
```

Output: 4 Parts

Goodness-of-fit for edgewise shared partner

	obs	min	mean	max	MC	p-value
esp0	3	2	11.33	19		0.04
esp1	9	0	2.53	12		0.10
esp2	3	0	0.25	6		0.06
esp3	0	0	0.02	1		1.00



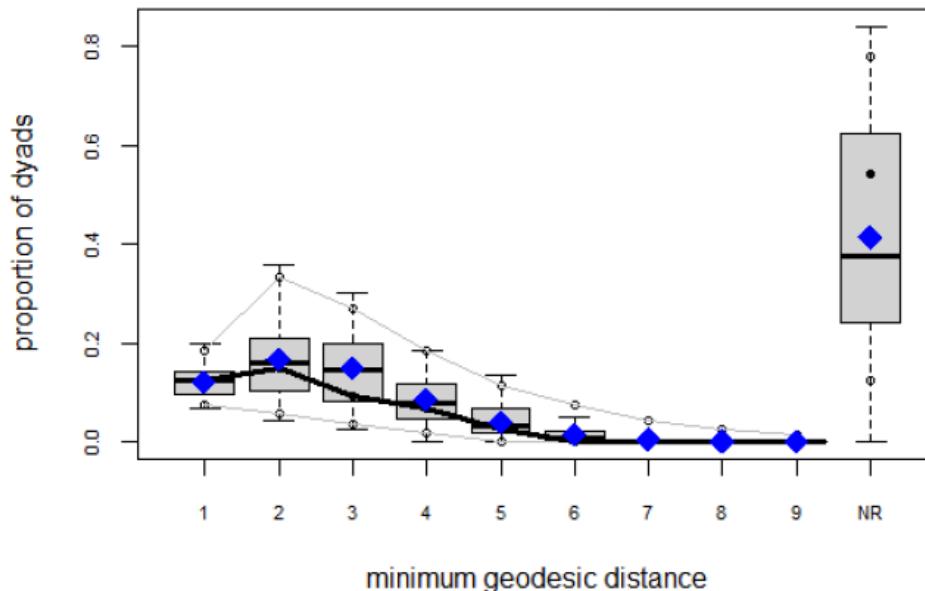
# GOF 3. Goodness-of-fit for minimum geodesic distance

```
fit <- ergm:::ergm(flobusiness ~ edges + degree(1))
fit.gof <- ergm:::gof(fit)
```

Output: 4 Parts

Goodness-of-fit for minimum geodesic distance

	obs	min	mean	max	MC	p-value
1	15	6	14.13	23		0.92
2	18	4	18.77	38		1.00
3	11	0	16.53	38		0.72
4	8	0	9.83	22		0.90
5	3	0	4.43	15		0.88
6	0	0	1.75	11		0.90
7	0	0	0.59	6		1.00
8	0	0	0.17	4		1.00
9	0	0	0.06	2		1.00
10	0	0	0.01	1		1.00
Inf	65	0	53.73	108		0.70
isolates						



## GOF 4. Goodness-of-fit for model statistics

```
fit <- ergm:::ergm(flobusiness ~ edges + degree(1))  
fit.gof <- ergm:::gof(fit)
```

Output: 4 Parts

### Goodness-of-fit for model statistics

	obs	min	mean	max	MC	p-value
edges	15	6	14.13	23		0.92
degree1	3	1	3.50	7		1.00

One box plot (estimate) for every stat in the model

